

Formulae

FRL-General

- If number of forks = n :

Total Number of processes created = 2^n

Number of child processes = $2^n - 1$

```
fork();  
fork();  
// ... n times  
printf("Hello!");
```

... will print("Hello!") 2^n times.

Basics

Goals of an Operating System

- Provide convenience to the user.
- Provide user-friendliness.
- Provide an interface to allow users to use applications to access & instruct the hardware.
- Hardware management
- Process management: Manage all currently running processes in the system. CPU Scheduling algorithms are used to determine which process will be executed by the CPU.
- Memory management: Manage volatile memory, dynamically allocate storage to ensure efficient utilization of the available volatile memory.
- Storage management: Manage non-volatile storage, using File System.

- Security:
 - Provide a certain level of security, so only authorized users can unlock and use the computer system.
 - Processes cannot access each other's data. Processes can only use the segment of CPU & RAM allocated to them.

Types of Operating Systems

Batch OS

- During 1960's, computers weren't so common. So we needed to go to a particular place, which provided computing services, to get the job done.
- The jobs were first loaded offline to a physical storage device like punch card, paper tape, magnetic tape, etc.
- They were then submitted to the operator.
- The Operator sorted jobs into batches.
- The first batch was provided to the CPU for execution. All jobs were executed by the CPU one by one.
- Since only 1 job was executed at a time, if the job needed I/O time, the CPU would remain idle in that time. This was a major disadvantage.
- When the CPU had finished executing the job and produced the result, it was loaded in the physical storage again, and given back to the user.
- Later on, IBM launched FORTRAN & IBSYS709X, which provided monitors where the user could directly punch the punch card.

Multi-programming

- The objective is to bring as many processes to the volatile memory as possible.
- It's nature is **non-preemptive**, ie only 1 process is executed by the CPU at any time. If the process needs I/O, the processor will remain idle during that time.
- The CPU won't move to the next process unless the current process has either finished executing or until the process tells it to move on, by itself.

Multitasking / Time Sharing

- It's nature is **pre-emptive**, ie each process is allocated a specific time interval to execute itself, regardless of how much time it needs. After the time-quantum expires, it has to leave the CPU to make room for another process.
- It ensures that no process has to wait for a very long amount of time to execute itself. All processes are executed within a reasonable amount of time, and no process is left out.
- It results in more responsiveness, compared to a Multi-programming OS.

Process States

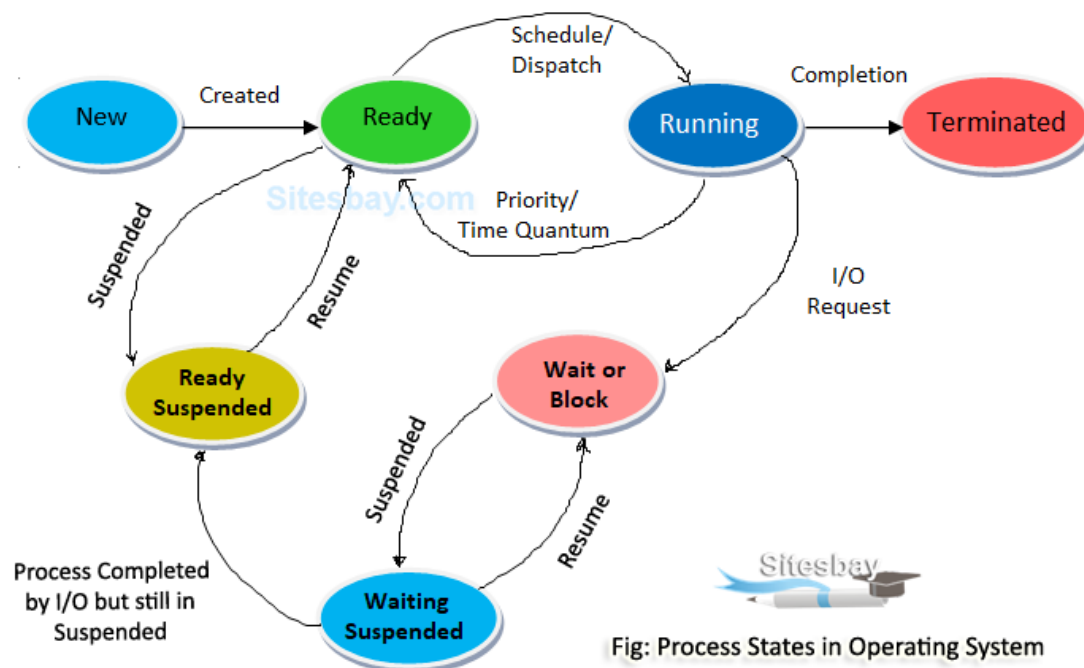


Fig: Process States in Operating System

Image taken from [here](#)

- There are **5 main states** & 2 suspend states: **New**, **Ready**, Suspend Ready, **Running**, **Wait/Block**, Suspend Wait, **Terminated**.
- New <-> **Long-Term Scheduler (LTS)** <-> Ready
- Ready <-> **Short-Term Scheduler (STS)** <-> Running
- Ready / Wait <-> **Medium-Term Scheduler (MTS)** <-> Suspend Ready / Suspend wait
- **New**: At first, a new process is created.

- **Ready:**

- The process is allocated memory & other resources as per needs.
- The **LTS (Long-Term Scheduler)** is responsible for bringing processes in from the `New` state.
- At the end of this state, it is ready to be executed.

- **Running, Wait:**

- The Job Scheduler queues all the processes in the `Ready` state, and dispatches it to the CPU for executing.
- The **STS (Short-Term Scheduler)** is responsible for moving processes back and forth the `Ready` & `Running` states.
- During execution, sometimes the process is moved back to `Ready` queue:
 - if a higher priority Process `HP` arrives in the `Ready` queue, the currently running process `P` is moved to the `Ready` state to make room for process `P`.
 - If the `Ready` queue is already full, the MTS (Medium-Term Scheduler) moves some processes to the `Suspend Ready` queue to make room for Process `HP`.
 - when using time-quantum based algorithms like Round Robin. If the time-quantum is 2 seconds, it means that each running process will be given 2 seconds of CPU Time, after which it'll have to make room for another process.
 - CPU initiates this process.
- During execution, sometimes the process requires some I/O operation. If that happens, it is moved to the `Wait/Block` state, where it can perform the operation it needs.
 - The secondary storage is usually much much slower than the CPU. This is done to reduce CPU idle time.
 - The process initiates this request.
 - After the I/O operation has finished, the CPU moves back to the `Ready` state.
- If a lot of processes need to execute I/O operations at the same time, the I/O queue will get full. In such a case, some processes are swapped out to the secondary memory ie moved to the `suspend wait` queue.
 - The **MTS (Medium-Term Scheduler)** is responsible for swapping

processes to and from `Wait/Block` state and `Suspend Wait` state.

- After the I/O operation is complete and the process is in the `Suspend Wait` state, it tries to get back to the `Wait` state. If the `Wait` queue is full for a significant amount of time, it is moved to the `Ready` queue. This is called `Backing Store`.
- **Terminated:** After the execution of the process is complete, all resources are de-allocated, and the process moves to the `Terminated` state.

System Calls

- By default, we use software/apps in **user-mode**.
- In user-mode, we don't have the rights to directly interact with the hardware. The kernel manages that part, as part of the operating system.
- When we want to communicate with or send information to hardware, we need to interact with the kernel, using `system calls`.
- So, a system call is used to invoke the kernel to perform operations on hardware, files, etc.

Types of System Calls

- **File-related:** During execution, if the process needs access to a particular file, it requests the kernel to provide it access, using file-related system calls. Example: `open()` , `read()` , `write()` , `close()` , `create file`.
- **Device-related:** We obtain the rights to access and talk to hardware, using these system calls. Example: `read`, `write`, `reposition`, `ioctl`.
- **Information-related:** We use these system calls when we want to get information about something. Example: `getPid` , `attributes`, `system-wide time & date`.
- **Process control:** These are used for managing processes. Example: `load()` , `execute()` , `abort()` , `fork()` , `nbit` , `signal` , `allocate` . `wait` & `signal` are used for process synchronization.
- **Communication:** These are used by processes to communicate among themselves. Example: `pipe()` , `create/delete connections`, `shmget()` [get value of the shared memory].
- **Security:** Here, we're mostly concerned with security & permissions. Examples: `chmod` , `chown` , `umask` .

chmod (Change mode)

```
$ ls -lh # Easiest way to display current permissions
total 12K
drwxr-xr-x 2 sayan sayan 4.0K Aug 13 20:39 dir1:
-          Directory?
---      Read, write, execute permissions for Owner / User who owns the fi
---      Read, write, execute permissions for Group
---      Read, write, execute permissions for Others (Everyone else)
```

Changing permissions:

- Can be done in 2 ways:

- Method 1:

- Permissions can either be defined specifically for `u/g/o` , or they can be combined together.
- `o-x` means we're stripping the `execute` permission off `others` .
- `ug+w` means we're providing `user` & `group` with the `write` permission.
- `a+x` means we're giving everyone (ie user, group & others) the `execute` permission.
- Example:

```
$ chmod o-x dir1 # Remove 'execute' permission for 'others'.
$ chmod o+w dir1 # Add 'write' permission for 'others'.
$ chmod a+x dir1 # Add 'write' permission for everyone (u,g as
$ chmod -R o+w dir1 # Add 'write' permission for 'others' recu.
$ chown abcd dir1 # Transfer directory ownership to user 'abcd'
$ chown -R abcd dir1 # Transfer directory ownership to user 'ab
```

- Method 2 (octal): `chmod <user><group><other> file` , or `chmod ugo file` .

- Meaning of the numbers: $r = 4$, $w = 2$, $x = 1$
- Permissions are denoted by numbers like 1, 2, 4 or a sum of any of the numbers, like 5, 6, 7.
- Permissions are defined in sequence, for `u` , `g` and `o` .
- Example:

```
chmod 111 abcd # u=x | g=x | o=x
chmod 666 abcd # u=rw | g=rw | o=rw
chmod 421 abcd # u=r | g=w | o=x
```

lseek()

- It is a system call that is used to change the location of the read/write pointer of the file descriptor.
- By default, pointer stays in the beginning, at index 0.
- Syntax & example (input file = 1234567890abcdefghij):

```
lseek(int file_descriptor, offset, int whence)
lseek(n,10,SEEK_CUR) # It goes 10 bytes from character `1`, ie to `a`.
lseek(n,5,SEEK_SET) # pointer is set at the position 5, ie at `5`.
```

- `file_descriptor` : The file descriptor of the pointer that is going to be moved.
- `offset` : The off-set of the pointer.
- `whence` : The method in which the offset will be interpreted. Possible values:
 - `SEEK_SET` : Set the off-set to the specified index.
 - `SEEK_CUR` : Off-set from current location of the pointer.
 - `SEEK_END` : Off-set from the end.

fork()

- It is used to create a child process, which is a clone of the parent process and has it's own PID.
- `fork()` returns different values depending on which process we're in:

```
0, if we're in the child process
+1/+ve number, ie the PID of the child process, if we're in the parent process.
-1/-ve number, if the child process couldn't be created.
```

- If `fork()` is run n times, it will create 2^n total processes, including $2^n - 1$ child processes and 1 parent process.
- Child process runs parallelly with the parent process.
- Example 0:

```
main() {
```

```

fork();
printf("Hello");
}

```


... will print "Hello" 2 times.

- Example 1:

```

#include<stdio.h>
#include<unistd.h>
int main() {
    if(fork() && fork()) { // 0, 1
        fork();           // 2
    }
    printf("Hello");
    return 0;
}

```




- Here, at fork() #0, a child process c_1 is created.

In the child process, it returns 0. The loop condition becomes false, and it exits.

In the parent process, it returns the PID of the child process. We proceed to fork() #1 .

- In fork() #1 , child process c_2 is created, which also returns 0. The loop exits. In the parent process, we go ahead and execute fork() #2 . This forks another child process c_3 .
 - In total, there are 4 processes, and `Hello` is printed 4 times.
- Example 2:

```
#include<stdio.h>
#include<unistd.h>
int main() {
    if(fork() || fork()) { // 0, 1
        fork();           // 2
    }
    printf("Hello");
    return 0;
}
```



- Here, at `fork() #0` :

In the child process, it returns 0.

In the parent process, it returns the PID of the child process.
- For the child process c_1 , it will continue to check the 2nd sub-condition, which is `fork() #1`. That will fork another child process c_4 .
- For the parent process, since the 1st condition is true, it won't even check the 2nd condition. It'll directly enter the loop, and execute `fork() #2` in it. That will fork another child process, c_2 .
- Now, c_1 becomes the parent process. Since the condition is still true, it gets in the loop and executes `fork()` in it. This creates another child process, c_4 .
- In total, there are 5 processes, and `Hello` is printed 5 times.