

FRL Definitions

- **State-Space:** The set of all possible states.
- **Optimal Solution:** A solution that meets certain criteria, such as minimizing cost, maximizing utility, or achieving the best possible outcome, given a specific problem and its constraints. It is the most desirable solution among all possible solutions available for a given problem.
- **Polynomial problems:** These are problems for which the time complexity of the best-known algorithm is polynomial in the size of the input. In other words, the time taken to solve these problems grows at most polynomially with the size of the input.
- **Non-polynomial problems:** These are problems for which the time complexity of the best-known algorithm is non-polynomial in the size of the input. In other words, the time taken to solve these problems grows exponentially or faster with the size of the input.
- **Heuristic:** The term **Heuristic** originates from the Greek word "heuriskein," which means "to discover" or "to find." It is a strategy that helps to efficiently navigate complex situations or tasks, often by simplifying the problem or focusing attention on relevant information.

Basics

- With Artificial Intelligence, our goal is to build machines and software with intelligence similar to humans. These machines will be able to perform thinking, reasoning, decision-making, problem solving, natural language processing, just like humans.
- The main goals of an artificially intelligent system are:
 - Reasoning
 - Learning
 - Problem Solving
 - Perception ... like a human.
- When it comes to solving a problem, you need to represent the problem in such a way that the machine can understand it. Represent:
 - Precisely
 - In such a way that it can be analyzed
- The major branches of AI are:
 - Perceptive
 - Vision
 - Robotics
 - Expert Systems: stores knowledge, makes inferences from it
 - Learning systems
- Views of AI:

Thinking humanly	Thinking Rationally
Acting humanly	Acting Rationally

- Things AI Can Do:
 - **Natural Language Processing (NLP):** AI can understand, interpret, and generate human language, enabling applications like chatbots and language translation.
 - **Image and Video Recognition:** AI can analyze and identify objects, faces, and scenes in images and videos, useful in security and medical imaging.
 - **Predictive Analytics:** AI can analyze historical data to make predictions about future events, aiding decision-making in finance and healthcare.
 - **Robotics and Automation:** AI can control robots and automate tasks, advancing manufacturing and logistics.
 - **Recommendation Systems:** AI can personalize content and product recommendations based on user behavior, enhancing experiences on e-commerce sites and streaming services.
- Categories of Artificially Intelligent Systems:
 - **Narrow AI (Weak AI):** These are AI systems designed to handle a specific task or a narrow range of tasks. Examples include virtual assistants like Siri and Alexa, recommendation algorithms on Netflix, and image recognition systems.
 - **General AI (Strong AI):** This is a theoretical concept where AI possesses the ability to understand, learn, and apply knowledge across a wide range of tasks, much like a human being. General AI can reason, solve problems, and make decisions autonomously in any situation.
 - **Superintelligence:** This refers to AI that surpasses human intelligence and capabilities. It can perform tasks and solve problems beyond human comprehension, potentially leading to breakthroughs in various fields. Superintelligence is still a hypothetical concept and has not been achieved yet.

Searching Techniques

- Uninformed & Informed Search
- Difference: **+**: Informed, **-**: Uninformed

```
+ Utilizes specific information about the problem domain to guide the search process.
- Lacks specific domain knowledge and relies solely on general search strategies. It is only aware of the start & goal state.
+ Generally more efficient in terms of time and space complexity.
- May be less efficient compared to informed search algorithms, especially for complex problems.
+ Designed to provide an optimal solution.
- May or may not provide an optimal solution.
+ Examples include A* search, heuristic search, and informed hill climbing.
- Examples include depth-first search, breadth-first search, and uniform-cost search.
```

State-Space Search


- State-Space: The set of all possible states.
- Set: $\{S, A, \text{Action}(s), \text{Result}(s,a), \text{Cost}(s,a)\}$

- S: Start, Goal
- A: The set of all possible actions.
- Action(s): The action we chose to execute.
- Result(s,a): State formed as a Result of the action.
- Cost(s,a): Cost of execute the action. The goal is to minimize the cost.

Breadth-First Search

- Type: Uninformed Search.
- Based on: FIFO (Queue).
- Time complexity: $O(b^d)$
 - b: Branch factor, maximum number of children of a node.
 - d: Depth: Maximum Level of the tree, root node is at Level 0.
- Optimal, provides the best solution, if costs of all nodes is the same.
- Complete, always provides a solution.
- Example 0 (Start: A, Goal: G):
 1. A
 2. ~~A~~BCD
 3. ~~B~~CDEF
 4. ~~C~~DEFGH
 5. ~~D~~EFGHI
 6. ~~E~~FGHIJK
 7. ~~F~~GHIJK
 8. ~~G~~Hijkl
 - G was found. Result: ACG.

- Note the implementation of FIFO: Elements are removed from LHS, and inserted from RHS.




Depth-First Search

- Type: Uninformed Search.
- Based on: LIFO (Stack).
- Time complexity: $O(b^d)$
 - b : Branch factor, maximum number of children of a node.
 - d : Depth: Maximum Level of the tree, root node is at Level 0.
- Not Optimal, may not provide the best solution.
- Not Complete, may not provide a solution.
- Example 0 (Start: A, Goal: D):
 1. A
 2. ~~A~~BC
 3. ~~B~~CFG
 4. B~~F~~G
 5. BF
 6. ~~B~~DE
 7. DE


8. \emptyset

- D was found. Result: **ACG**
- Sequence: **ACGFBED**
- Note the implementation of LIFO: Elements are removed from RHS, and inserted from RHS.



Bi-directional Search

- Type: Depends on algorithm used.
- 2 simultaneous search, one from initial node to goal node, another from goal node to initial node.
- Time complexity: $O(b^d + b^d) = O(2b^{d/2})$
 - b: Branch factor, maximum number of children of a node.
 - d: Depth: Maximum Level of the tree, root node is at Level 0.
- Complete only in case of Breadth-First Search.



- Type: Blind / Uninformed Search.
- Based on: Breadth-First Search
- Time complexity: $O(b^d)$
 - b: Branch factor, maximum number of children of a node.
 - d: Depth: Maximum Level of the tree.
- Example 0:
 - Actions (A): UP (↑), DOWN (↓), LEFT (←), RIGHT (→)
 - Start | End:

S				G		
1	2	3	..	1	2	3
				4	5	6
7	5	8		7	8	

1. Step 1:

0				->	↑			→				↓		
1	2	3	->		2	3	..	1	2	3	..	1	2	3
4	6	->		1	4	6		4		6		7	4	6
7	5	8	->	7	5	8		7	5	8		5	8	

2. Step 2 (from puzzle →):

0			->	↑			↓			←			→					
1	2	3	->	1		3	..	1	2	3	..	1	2	3	..	1	2	3
4		6	->	4	2	6		4	5	6		4	6		4	6		
7	5	8	->	7	5	8		7		8		7	5	8		7	5	8

3. Step 3 (from puzzle ↓):

0			->	←			↑			→				
1	2	3	->	1	2	3	..	1	2	3	..	1	2	3
4	5	6	->	4	5	6		4		6		4	5	6
7		8	->		7	8		7	5	8		7	8	

4. Step 4: Puzzle 3 is the Goal State.

- At every step, all valid moves are executed for all states, and all resultant states are determined.

Heuristic in Artificial Intelligence

- The term "heuristic" originates from the Greek word "heuriskein," which means "to discover" or "to find." It is a strategy that helps to efficiently navigate complex situations or tasks, often by simplifying the problem or focusing attention on relevant information.
- We use heuristic functions when we want to convert non-polynomial problems to polynomial problems (NP→P).

- **Polynomial problems:** These are problems for which the time complexity of the best-known algorithm is polynomial in the size of the input. In other words, the time taken to solve these problems grows at most polynomially with the size of the input.
- **Non-polynomial problems:** These are problems for which the time complexity of the best-known algorithm is non-polynomial in the size of the input. In other words, the time taken to solve these problems grows exponentially or faster with the size of the input.
- It provides a good solution but **not an optimal solution**. This is because there may be other obstacles (like an infinite loop) in the path the algorithm has chosen. But in most cases, it provides a more efficient approach to brute-forcing our way to the solution.
- Examples of some functions:
 - **Eucledian Distance:** $\sqrt{(x_2-x_1)^2+(y_2-y_1)^2}$, using this we can find the shortest path from point a to point b.
 - **Manhattan Distance:**

S			->	G		
1	3	2	->	1	2	3
6	5	4	->	4	5	6
8	7		->	7	8	

- Manhattan Distance: $0+1+1+2+0+2+2+0=8$
- The distance is calculated by the number of steps each number needs to be moved, to reach the Goal state (G) from the Current state.

Searching Techniques

8-Puzzle Problem with Heuristic

- Type: Informed Search.
- Based on: Number of misplaced tiles
- Example 0:
 - Actions (A): UP (↑), DOWN (↓), LEFT (←), RIGHT (→)
 - Start | End:

S				G			
1	2	3	..	1	2	3	
	4	6		4	5	6	
7	5	8		7	8		

- Number of misplaced tiles, $h=3$

1. Step 1:

0			->	↑			→			↓				
1	2	3	->		2	3	..	1	2	3	..	1	2	3
	4	6	->	1	4	6		4		6		7	4	6
7	5	8	->	7	5	8		7	5	8			5	8

- Number of misplaced tiles, $h=4,2,4$

- We will move forward with the lowest heuristic value.
2. Step 2 (from puzzle ➡):

0			->	↑			↓			←			➡					
1	2	3	->	1		3	..	1	2	3	..	1	2	3	..	1	2	3
4		6	->	4	2	6		4	5	6			4	6		4	6	
7	5	8	->	7	5	8		7		8		7	5	8		7	5	8

- Number of misplaced tiles, \$h=3,1,3,3\$
3. Step 3 (from puzzle ↓):

0			->	←				↑				➡		
1	2	3	->	1	2	3	..	1	2	3	..	1	2	3
4	5	6	->	4	5	6		4		6		4	5	6
7		8	->		7	8		7	5	8		7	8	

- Number of misplaced tiles, \$h=2,2,0\$
4. Step 4: For Puzzle ➡, \$h=0\$, and it is the Goal State.
5. We had to traverse through a lot less states to reach the Goal state (G), compared to a typical Uninformed Search Technique.

Generate and Test

1. Generate a possible solution.
2. Test to see if this is an actual solution.
3. If a solution is found, otherwise repeat.

- Properties of Good Generators:
 - **Complete**
 - **Non-redundant**: They must not provide solutions which have already been generated in the past.
 - **Informed**: The Generator must have atleast some basic idea which it can use to generate an efficient solution.

Machine Learning

- Artificial Intelligence: Aimed at enabling computers to perform human-like tasks and simulate human behaviour.
- Machine Learning: Tries to solve a specific problem, and makes predictions using data.
- Data Science: Attempts to find patterns and find insights from data.

Types of Machine Learning

- Supervised: Learning where the model is trained on labeled data (all tuples have a class label associated with them) and learns to make predictions based on examples.
- Unsupervised: Learning from data without labeled responses, finding patterns and relationships on its own.

- Reinforcement: Learning through trial and error, where an agent learns to make decisions by interacting with an environment and receiving feedback (either positive/reward or negative/penalty).