

Questions

qs-UGC_NET-2018

Let $R_1(a, b, c)$ and $R_2(x, y, z)$ be 2 relations in which 'a' is a foreign key in R_1 that refers to the primary key (unspecified) of R_2 . Consider 4 options:

1. Insert into R_1
2. Insert into R_2
3. Delete from R_1
4. Delete from R_2

Which of these options is true?

- Option 'a' & 'c' will cause violation.
- Option 'b' & 'c' will cause violation.
- Option 'c' & 'd' will cause violation.
- **Option 'd' & 'a' will cause violation.**

Explanation: Check

► Explanation

qs-GATE-2005

What is the minimum number of tables required to represent this E-R Model into Relational Model?

- 2
- **3**
- 4
- 5

Basics

Structured data

- Data that needs to be stored in a structured manner.
- Data stored in a RDB, ie a Relational Database, in the form of a relation/table. It is managed by a RDBMS.
- Users can operate on the data using DBMS/RDBMS. Possible operations: Insert, Add, Delete, Update.
- Examples: SQL Server, Oracle, MySQL, etc.

Unstructured data:

- Data that does not need to be stored in a particular pre-defined structure.
- Example: Websites, photos, videos, etc.

File System vs DBMS

- Filesystems were used before DBMS existed.
- In 1970's, the user used to access data in their own system.
- Advantages of using databases & database management systems (- : filesystem, + : database):
 - Data is stored in 1 user's computer, and is only accessible by that person.
 - + We have the data stored in a centralized server, and multiple users can access it.
 - Users need to download the full file even if they only want some specific data.
 - + Users can send a query to the server to access only the data they need.
 - Users need to know the location of the file, to fetch it.
 - + Users don't need to know details about the data ie where it is stored.
 - Concurrency: In modern times, a lot of people access/update data from their own system.
 - + There are protocols in place to handle inconsistency & conflict during concurrent access.

- Security: Since filesystem is controlled by OS, there is no level-by-level security.
- + Role-based (hierarchical) security protocols are available, so different users can have different access.
- Redundancy: Multiple files with same content but different names, can be stored.
- + DBMS has various constraints to ensure we don't waste storage space by storing redundant data.

2-Tier Architecture

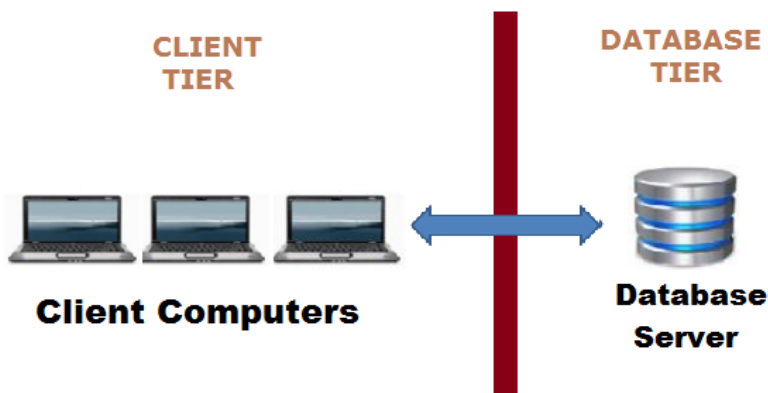


Image taken from [here](#)

- Consists of multiple clients, directly interacting with a single database server.
- Clients have an interface which help them interact with the database server.
- Client connects to the database, sends the query, and fetches the information.
Information can also be modified or deleted from the database server as needed.

- + Simple architecture.
- + Easy to maintain.
- Difficult to scale.
- Implementing security is difficult, because client is directly interacting with the database server.

3-Tier Architecture:

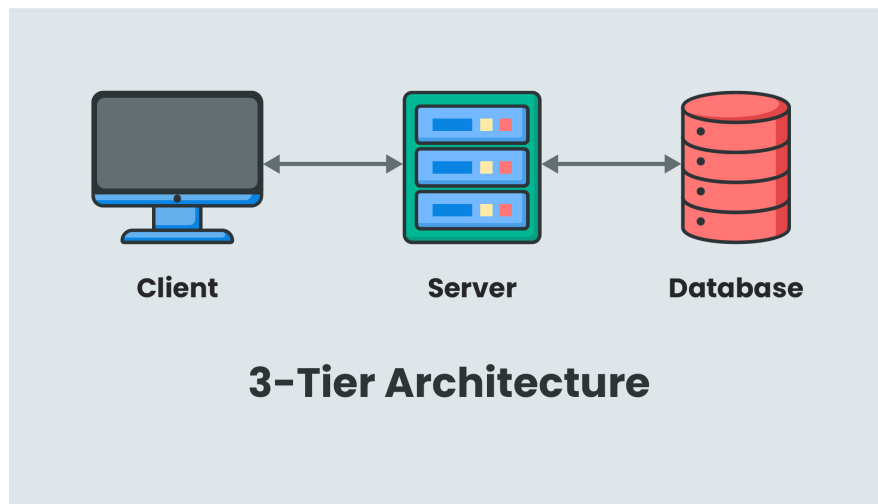


Image taken from [here](#)

- 3 Layers: Client (Application) Layer > Application Server / Business Layer > Database Layer
- Clients have an interface which allows them to interact with the middleware application server. They don't directly connect to the database server.
- Database server does not get overwhelmed since the middleware server handles all the queries, so load on it is reduced.

+ Can be scaled.

+ Secure, since client is not directly interacting with the database server

- More complex and difficult to maintain compared to 2-Tier architecture.

Schema & 3-Schema Architecture

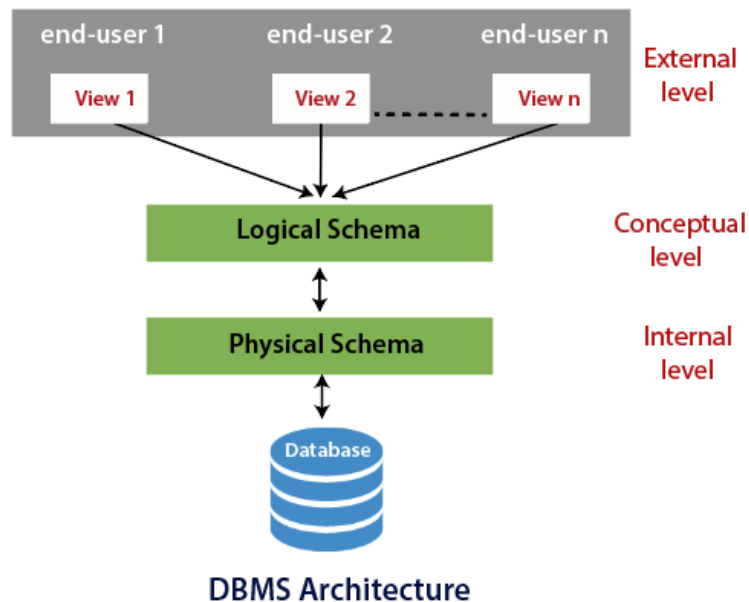


Image taken from [here](#)

- Schema means the logical representation of data in the database. It tells us about the structure in which the data will be stored. We also get information about the elements in the structure, their size, etc.
- We're not interested in the physical representation of the data, ie how exactly they're stored in the database. We're only interested in a legible representation of the data.
- A 3-Schema architecture consists of 3 levels: View | Logical | Physical
- We see data in the form of a table, but it is actually stored in the form of files.

1: External Schema / View Level:

- Controlled by the Front-end Developer.
- Different types of users (student, faculty, etc.) have different views.
- Within a particular type, different users are shown different data, relevant to them, within the particular view.

2: Conceptual Schema / Logical Level:

- Controlled by the Database Designer.
- Can be called a blueprint of the structure.
- Stores information about the tables, and the data in the tables.
- Stores Relationships between the tables ie how they are connected.

3: Physical Schema / Physical Level:

- Controlled by the Database Administrator.
- Stores information about where the data is stored, how it is stored.
- Decides if data is centralized or de-centralized.

Data Independence

- The actual data is shown to the user, but information about how it is stored, where it is stored, etc. is hidden.
- We don't show all the tables to the user, we only show the relevant information the user wants to see.
- The application program is not changed, since the user directly interacts with it and it used to the particular view.
- The objective is to make changes in the Conceptual and Physical Schema in such a way that the user (View Level) is not affected.
- **Logical Data Independence:**
 - Changes made in the Conceptual Schema don't affect the External Schema.
 - Example: The WebUI shown to an user is not affected even after making some changes to the Conceptual Schema of a database system.
- **Physical Data Independence:**
 - Changes made in the Physical Schema don't affect the Conceptual Schema or the External Schema.
 - Example: Changing file structure, physical storage, Indexing, etc don't affect the Conceptual Schema.

Keys

- Any piece of data within a tuple in a database, can be called a key.
- Keys are of many types: Primary Key, Candidate Key, etc.

Primary & Candidate Key

- Any key that can be used to uniquely identify data tuples (or rows) in a database, is called a primary key.
- The primary key must be unique. Also, it must not be wrong or have any default value.
- It must not be null.
- There must only be 1 primary key in a database table (in case of RDB).
- Usually, if there's a possibility that the user cannot provide the data for a primary key at this time, it is auto-generated and assigned to the user. For example, registration number, roll number, etc.
- The set of all possible keys that can be used as a primary key, is called Candidate Key.
- We can have keys whose value is null, in the set of Candidate Keys. However, the `Unique` property must be maintained.
- Example:

Roll Number	Registration Number	SName	City	Age	Phone Number
101	REG2022001	Alice	Paris	20	555-1234-5678
102	REG2022002	Bob	London	22	555-1234-5678
103	REG2022003	Charlie	Berlin	21	555-3456-7890
104	REG2022004	David	Tokyo	23	null
105	REG2022005	Emily	Rome	20	555-5678-9012

- Here, the candidate keys are: <Roll No, Registration Number, Phone Number>
- Among these, Phone number can be null, so it can't be used as a primary key. Any 1 of the other 2 can be used for the purpose.

Foreign Key

- An attribute or set of attributes that references to `Primary Key` of the same table or another table.
- It maintains referential integrity.
- If 2 tables are related, there will be atleast 1 common column between them. However, it's not necessary for this column to have the same name in both the tables.
- Example: **Table 1:** Student Information

Roll No	Name	Address
101	Alice	Paris
102	Bob	London
103	Charlie	Berlin
104	David	Tokyo

Table 2: Course Information

Course ID	Course Name	Roll Number
201	Math	101
202	Science	102
203	English	103
204	History	104

- Here, `Roll Number` of Table 2 is a foreign key, referencing `Roll No` in Table 1.
- The names of the columns may or may not be unique.
- SQL query for setting a foreign key during table creation:

```
CREATE TABLE StudentInformation (
```



```

RollNo INT PRIMARY KEY,
Name VARCHAR(20),
Address VARCHAR(20)
);

CREATE TABLE CourseInformation (
    CourseID INT PRIMARY KEY,
    CourseName VARCHAR(20),
    RollNumber INT,
    FOREIGN KEY (RollNumber) REFERENCES StudentInformation(RollNo)
);

```

- SQL query for setting a foreign key after table creation (CONSTRAINT name can be anything as it's the name of the constraint):

```

ALTER TABLE CourseInformation
    ADD CONSTRAINT FK_StudentInfo_RollNo
    FOREIGN KEY (RollNumber) REFERENCES StudentInformation(RollNo);

```

Referential Integrity

Integrity means consistency of the data across multiple tables. By preserving referential integrity, we aim to ensure consistency among the data across both the referenced (base) & referencing table.

- Now, let's check if INSERT, DELETE & UPDATE operations on either of the tables disturbs the Referential Integrity of the data.
- Example: **Referenced (Base) Table 1: Student Information**

Roll No	Name	Address
101	Alice	Paris
102	Bob	London
103	Charlie	Berlin
104	David	Tokyo

- Primary Key: Roll No

Referencing Table 2: Course Information

Course ID	Course Name	Roll Number
201	Math	101
202	Science	102
203	English	103
204	History	104

- Foreign Key: Roll Number REFERENCES Roll No .
- Operations performed on Referenced (Base) Table 1 :
 - **INSERT**: No violation
 - **UPDATE**: Possible violation.
 - Updating the Roll No on Table 1 can introduce inconsistency, because the tuple in Table 2 that was referencing to it, now links to nothing.
 - Solution:
 - i. Update manually: Recommended.
 - ii. *onUpdateCascade()*: When Roll No is updated, delete the corresponding row on Table 2 . This can result in loss of data.
 - iii. *onUpdateSetNull()*: When Roll No is updated, set the Roll Number of the corresponding row on Table 2 , to null. However, this may cause issues if the foreign key is also a primary key of Table 2 .
 - **DELETE**: Possible violation.
 - We may accidentally delete a tuple whose Roll No has been referenced by Table 2 .
 - Solution:
 - i. No action: Do nothing, this is not recommended for obvious reasons.
 - ii. *onDeleteSetNull()*: When the tuple is deleted, set the Roll Number of the corresponding row on Table 2 , to null .
- Operations performed on Referencing Table 2 :
 - **INSERT**: Possible violation.

- We cannot insert a tuple in this table with a Roll Number that doesn't reference to any tuple in Table 1 .
- Solution: Ensure that the corresponding tuple exists in Table 1 .
- **UPDATE:** Possible violation.
 - Updating the foreign key, Roll Number on Table 2 can introduce inconsistency, while it's fine to update other columns.
 - Solution:
 - i. Before updating the foreign key, ensure that it's not referencing to any primary key in Table 1 .
- **DELETE:** No violation.

Super Key

- It is a combination of all possible attributes which can uniquely identify 2 tuples in a table. We pair the candidate key with any other key, to uniquely identify the tuple.
- A super-set of any candidate key is called a super key.
- Example 0: IF $R(A_1, A_2, A_3, \dots, A_n)$:
 - If A_1 is a candidate key, then how many super keys are possible?
 - Usually, per key, we have 2 choices. We can either include it in our set, or exclude it.

Total number of possible choices in general: 2^n
 - But, here A_1 is said to be the candidate key, and must be included. For the rest, they may or may not be included.

Total number of possible choices: 2^{n-1}
 - If A_1 & A_2 are candidate keys:
 - If we just take A_1 , we have 2^{n-1} possible choices.
 - If we just take A_2 , we have 2^{n-1} possible choices.
 - **However**, both these super sets will have some elements in common. For A_1 set, we will have elements which contain A_2 in them, and vice versa. So, we need to exclude them. We intend to exclude elements which contain both A_1 & A_2 in them, which is 2^{n-2} .

Number of possible sets: $2^{n-1} + 2^{n-1} - 2^{n-2}$

- If A_1A_2 & A_3A_4 are candidate keys:

 | Similarly, the number of possible sets: $2^{n-2} + 2^{n-2} - 2^{n-4}$

Data Models

- Before implementing a database, we need to draw a logical design for the end-product. This design is called a data model.
-

Entity Relationship Model

- Entity: Object which has some attributes.
- Relationship: Relationships among the entities.

Attributes

Example: Student

1. Single: Attribute which has only 1 value.

 | Registration number.

2. Multivalued: has more than 1 values.

 | Mobile Number

1. Simple: cannot be broken further

 | Age

2. Composite: can be broken further

 | Name, can be broken down into first name, middle name, last name
 | Date of Birth, can be broken down into DD,MM,YY
 | Address

3. Complex: Composite + Multivalued

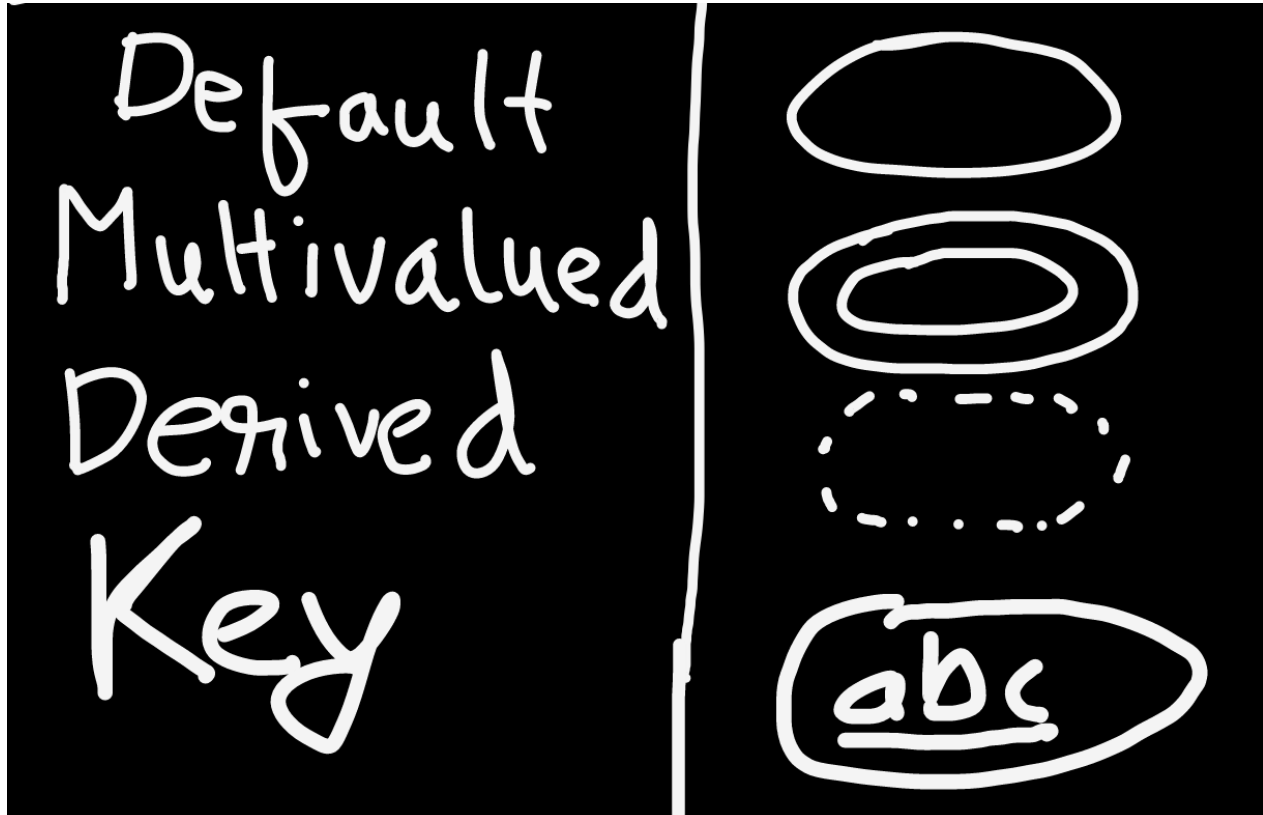
 | Address, it is composite, and a single student can have multiple addresses too.

1. Stored: directly stored.

| Date of Birth

2. Derived: derived from stored attribute.

| Age, derived from Date of Birth.



1. Key: can be used as the Primary key in a database.

| Registration number

2. Non-key: All keys other than the key attribute.

| Name

1. Required: mandatorily needs to have a value.

| Username, usually.

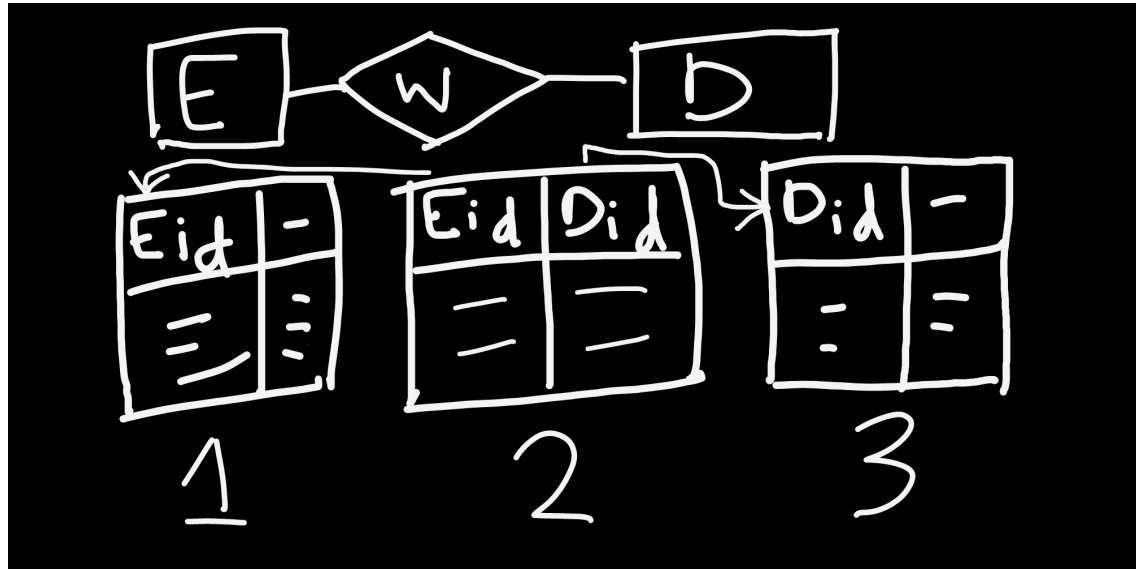
2. Optional: may not have a value.

| Mobile number, usually.

Relationships

One to One

- Representation: 1-1 | One-One
- Example:
 - 3 tables: Employee , Work , Department



- In Employee : E_id is the primary key
- In Department : D_id is the primary key
- In the Relationship table, Work :
 - Foreign keys:
 - E_id: referencing to E_id of Employee .
 - D_id: referencing to D_id of Department .
 - Relationship (1-1): An employee only works in a single Department.
 - E_id: Value is unique
 - D_id: Value is unique
 - Primary Key: Either of the primary keys of the LHS or RHS table.
 - Any of E_id or D_id
- We can reduce (merge) Employee & Work into a single Employee table. In fact, we can merge the Relationship table with either the LHS or the RHS table.

One to Many

- Representation: 1-M | One-Many
- Example:
 - 3 tables: Customer , Places , Order
 - In Customer : C_id is the primary key

- In `Order` : `O_no` is the primary key
- In the Relationship table, `Places` :
 - Foreign keys:
 - `C_id`: referencing to `C_id` of `Customer` .
 - `O_no`: referencing to `O_no` of `Order` .
 - Relationship (1-Many): A customer can place multiple orders.
 - `C_id`: Value can repeat
 - `O_no`: Value is unique
 - **Primary Key**: Primary key of the RHS table.
 - `O_no` , since `C_id` can repeat.
- We **can merge** (reduce) `Places` & `Order` into a single `Order` table, ie we're merging the Relationship table with the RHS table.

Many to One

- Representation: M-1 | Many-One
- Example:
 - 3 tables: `Order` , `Placed-by` , `Customer`
 - In `Order` : `O_no` is the primary key
 - In `Customer` : `C_id` is the primary key
 - In the Relationship table, `Placed-by` :
 - Foreign keys:
 - `O_no`: referencing to `O_no` of `Order` .
 - `C_id`: referencing to `C_id` of `Customer` .
 - Relationship (Many-1): Many orders can be placed by a single customer.
 - `O_no`: Value is unique
 - `C_id`: Value can repeat
 - **Primary Key**: Primary key of the LHS table.
 - `O_no` , since `C_id` can repeat.
 - We **can merge** (reduce) `Placed-by` & `Order` into a single `Order` table, ie we're merging with the LHS table.

Many to Many

- Representation: M-N | Many-Many
- Example:
 - 3 tables: `Student` , `Study` , `Course`
 - In `Study` : `RollNo` is the primary key.
 - In `Course` : `Cid` is the primary key.
 - In the Relationship table, `Study` :
 - Foreign keys:
 - `RollNo`: referencing to `RollNo` of `Student` .
 - `C_id`: referencing to `C_id` of `Course` .
 - Relationship (May-Many): A student can take up multiple courses, and a course can be taken up by multiple students.
 - `RollNo` : Value can repeat
 - `C_id` : Value can repeat
 - **Primary Key**: A composite key, including both the primary keys from the LHS & RHS tables.
 - `RollNo + C_id` , since both `RollNo` & `C_id` can repeat.
 - We **cannot merge** (reduce) the tables.

Normalization

- It is a method of removing or reducing redundancy ie duplicate values from a table.
- 2 rows cannot be exactly the same, since we have a primary key in every table. At the very least, the primary key column will be different. However, the other values may be partially or fully duplicated.
- There are 3 types of anomalies:
 - Anomaly: Error that only occurs under special circumstances.
 - Insertion
 - Deletion
 - Updation
- Example:

S_{ID}	Sname	C_ID	Cname	F_ID	Fname	Salary
101	Alice	201	Math	301	John	50000
102	Bob	202	English	302	Emily	55000
103	Charlie	203	Science	303	Sarah	60000
104	Bob	202	English	302	Emily	55000
105	Emma	205	Geography	305	Alex	58000

- **Insertion Anamoly:** We cannot insert a new course in this table, unless there's also a student to take it up, and S_{ID} is a primary key.
- **Deletion Anamoly:** We can delete $S_{ID} = 102$ safely without losing data because details for $C_{ID} = 202$ & $F_{ID} = 302$ is also present in other rows. But, if we delete $S_{ID} = 105$, we may lose important information.
- **Updation Anamoly:** Say we want to update the salary of faculty $F_{ID} = 302$ ie Emily from 55000 to 60000. Now, since there is only 1 faculty with $F_{ID} = 302$, the updation should only happen once. But, in reality, since there are multiple entries of $F_{ID} = 302$, value 55000 will be updated multiple times within the table (twice, in this example).
- Normalization aims to prevent these anamolies from happening.

Functional Dependency

- $X \rightarrow Y$: means that X determines Y, or Y is determined by X.
 - X: Determinant
 - Y: Dependent Attribute
- Example:
 - X: Sid, Y: Sname, $X \rightarrow Y$

i. VALID, data has been duplicated.

Sid	Sname
1	Ranjit
1	Ranjit

ii. VALID, they are different people.

Sid	Sname
1	Ranjit
2	Ranjit

iii. VALID, they are different people.

Sid	Sname
1	Ranjit
2	Varun

iv. INVALID, Sid cannot be same for different people.

Sid	Sname
1	Ranjit
1	Varun

- Types:

- **Trivial:** Functional Dependencies where if $(X \rightarrow Y)$, then Y is a subset of X. Also, they are always valid, because effectively, X is determining itself. Example: $(Sname, Sid \rightarrow Sid)$.
- **Non-Trivial:** Functional Dependencies where if $(X \rightarrow Y)$, then Y is not a subset of X. Example: $(Eid \rightarrow Ename)$.

- Properties:

- **Reflexivity:** If X & Y are Trivial Functional Dependencies (Y is a subset of X), then $(X \rightarrow Y)$
- **Augmentation:** If $(X \rightarrow Y)$, then $(XZ \rightarrow YZ)$
- **Transitive:** If $(X \rightarrow Y)$ & $(Y \rightarrow Z)$, then $(X \rightarrow Z)$
- **Union:** If $(X \rightarrow Y)$ & $(X \rightarrow Z)$, then $(X \rightarrow YZ)$
- **Decomposition:** If $(X \rightarrow YZ)$, then $(X \rightarrow Y)$ & $(X \rightarrow Z)$. You cannot Decompose the left hand side.
- **Composition:** If $(X \rightarrow Y)$ & $(Z \rightarrow W)$, then $(XZ \rightarrow YW)$
- **Pseudotransitivity:** If $(X \rightarrow Y)$ & $(WY \rightarrow Z)$, then $(WX \rightarrow Z)$

Closure Method

- A method to find candidate keys from functional dependencies (FDs).
- 2 methods:
 - i. One by one: Examples 0,1
 - ii. More efficient: Example 2
- Example 0:
 - $R=(ABCD) \mid FD = \{A \rightarrow B, B \rightarrow C, C \rightarrow D\}$
 - $A^+ = ABCD$
 - $B^+ = BCD$
 - $C^+ = CD$
 - $D^+ = D$
 - $AB^+ = ABCD$, but cannot be considered because candidate keys should be minimal. However, it can be a [super key](#).
 - Prime Attributes: $\{A\}$
 - These attributes can be used as a candidate key.
 - Non-Prime Attributes: $\{B, C, D\}$
 - These attributes cannot be used as a candidate key.
 - A is included in A^+ because it will always determine its own value.
 - Only attributes that determine the values of all attributes in the relation, can be the candidate key.
- Example 1:
 - $R=(ABCD) \mid FD = \{A \rightarrow B, B \rightarrow C, C \rightarrow D, D \rightarrow A\}$
 - $A^+ = ABCD$
 - $B^+ = BCDA$
 - $C^+ = CDAB$
 - $D^+ = DABC$
 - Prime Attributes: $\{A, B, C, D\}$
 - Non-Prime Attributes: $\{null\}$
- Example 2:
 - $R=(ABCDE) \mid FD = \{A \rightarrow B, BC \rightarrow D, E \rightarrow C, D \rightarrow A\}$
 - i. Check the RHS of the functional dependencies: AB, D, C, A . E is missing, so assume that E will be a part of the candidate key.

ii. Start check from A :

$AE^+ = \{AECBD\}$. Candidate Key Set: $\{AE\}$

iii. Check if A or E is in RHS: $[D \rightarrow A]$. Check for D :

$DE^+ = \{DEABC\}$. Candidate Key Set: $\{AE, DE\}$

iv. Check if D is in RHS: $[BC \rightarrow D]$. Check for B & C :

$BE^+ = \{BECDA\}$. Candidate Key Set: $\{AE, DE, BE\}$

$CE^+ = \{CE\}$. Candidate Key Set: $\{AE, DE, BE\}$

- So, in short, when we find a new candidate key, we search if that new attribute is also present in RHS of the FD. If it is, we check for the corresponding LHS attributes.

First Normal form (1-NF)

- The table should not contain **multivalued attributes**.
- Example: Normalize this table to 1-NF (Solution 3 is the one that works):

RollNo	Name	Course
1	Sai	C,C++
2	Harsh	Java
3	Onkar	C,DBMS

- Solution 1: Put only 1 course per row.

RollNo	Name	Course
1	Sai	C
1	Sai	C++
2	Harsh	Java
3	Onkar	C
3	Onkar	DBMS

- Primary Key: RollNo+Course

- Data Duplication

o Solution 2: Divide the Course column:

RollNo	Name	Course 1	Course 2
1	Sai	C	C++
2	Harsh	Java	null
3	Onkar	C	DBMS

- Primary Key: RollNo

- null : no value

- If there are 2 columns for course & someone takes up a 3rd course

- For people who have only taken 1 course, we have to put the value

o Solution 3: Segregate into 2 tables.

- Base Table

RollNo	Name
1	Sai
2	Harsh
3	Onkar

- Primary Key: RollNo

- Referencing Table:

RollNo	Course
1	C
1	C++
2	Java
3	C
3	DBMS

- **Primary Key:** RollNo+Course
- **Foreign Key:** RollNo , referencing to RollNo .

Second Normal form (2-NF)

- Table must be in First Normal Form (1-NF).
- All the non-prime attributes must be fully functionally dependent on the Candidate Key. They must not be only partially dependent on a part of (ie a proper subset of) the Candidate Key.
- If $(AB \rightarrow C)$, AB is a candidate key (prime attribute) while C is a non-prime attribute. If $(A \rightarrow C)$ or $(B \rightarrow C)$ is true, then the table is not in 2-NF.
- Example:

Customer ID	Store ID	Location
1	1	Delhi
1	3	Mumbai
2	1	Delhi
3	2	Bangalore
4	3	Mumbai

- **Primary Key:** Customer ID + Store ID
- Here, Location (non-prime attribute), is only dependent on one of the prime attributes (Store ID). So, it's not in 2-NF .
- So, we divide it into 2 tables.
- Table 1 (no dependency problems):

Customer ID	Store ID
1	1
1	3
2	1
3	2
4	3

- Table 2 (Location is dependent on the only prime attribute, `Store ID`):

Store ID	Location
1	Delhi
2	Bangalore
3	Mumbai

- Example 1:
 - $R = \{A, B, C, D, E, F\}$, FD: $\{C \rightarrow F, E \rightarrow A, EC \rightarrow D, A \rightarrow B\}$
 - Here, C & E are not present in the RHS of the functional dependencies. So, let's check for CE^+ .
 - $CE^+ = CEAFDB$
 - Candidate Key: CE
 - None of C,E are present in the RHS of the functional dependencies, so {CE} is the only candidate key.
 - Prime attributes: $\{C, E\}$
 - Non-Prime attributes: $\{A, B, D, F\}$
 - Here, because $C \rightarrow F, E \rightarrow A$, we can clearly see the C & E by themselves are determining a non-prime attribute. So, table is not in 2-NF.

Third Normal form (3-NF)

- Table must be in Second Normal Form (2-NF).
- Table must not have any Transitive Dependency. A prime attribute can determine a non-prime attribute, but a non-prime attribute cannot.
- If $AB \rightarrow CD$, and $D \rightarrow A$:
 - Candidate Keys: $\{AB, DB\}$
 - Prime Attributes: $\{A, B, D\}$
 - Non-prime Attributes: $\{C\}$
 - Among the Functional Dependencies, none have any Transitive Dependency.

Boyce Codd Normal Form (BCNF)

- A special case of 3rd Normal Form.

- Table must be in Third Normal Form (3-NF).
- The left hand side of each functional dependency should be a candidate key or super key.
- 3-NF preserves dependencies, but BCNF does not.
- However, BCNF ensures Lossless Decomposition.
- Example 0:
 - Candidate Keys: {RollNo, VoterID}
 - Functional Dependencies:
 - RollNo \rightarrow Name
 - RollNo \rightarrow VoterID
 - VoterID \rightarrow Age
 - VoterID \rightarrow RollNo
 - All are valid, since the LHS attributes are all prime attributes.
- Example 1: R(ABCD)
 - AB \rightarrow CD
 - D \rightarrow A
 - i. Candidate Keys: {AB,DB}
 - ii. Prime Attributes: {A,B,D}
 - iii. Non-Prime Attributes: {C}
 - iv. Here, in {D \rightarrow A}, D is not a candidate key. So the table is **not** in BCNF. Let's convert it into BCNF.
 - v. We separate DA since {D \rightarrow A} was the problem.
 - vi. We also separate BCD.
 - vii. So, ABCD \rightarrow DA,BCD
 - viii. Lossless Decomposition: The attribute that is common to both tables will be the candidate key of the tables individually, or both of them.
 - ix. Now, generate the Functional Dependencies: D \rightarrow A, BD \rightarrow C.
 - x. Here, {AB \rightarrow CD} is not preserved. So, the fact that BCNF does not preserve dependencies, is proved.

Decomposition

- Lossy Decomposition: Decomposition that leads to inconsistency in data.
- When Decomposing a table, the common attribute should be the candidate key of

either of the tables, or both of them. This ensures Lossless Decomposition.

- Example 0 (Lossy Decomposition):

- Table R :

A	B	C
1	2	1
2	2	2
3	3	2

- Table R_1 :

A	B
1	2
2	2
3	3

- Table R_2 :

B	C
2	1
2	2
3	2

- Question: Find the value of C if $A = 1$.

- Query: Select

R_2C from R_2 NATURAL JOIN R_1 , WHERE $R_1A = 1$

i. First, find the cross product of all rows of the tables.

A	B	B	C
1	2	2	1
1	2	2	2
1	2	3	2
2	2	2	1

A	B	B	C
2	2	2	2
2	2	3	2
3	3	2	1
3	3	2	2
3	3	3	2

ii. Remove the rows where the value of B is different.

A	B	B	C
1	2	2	1
1	2	2	2
2	2	2	1
2	2	2	2
3	3	3	2

iii. The Natural Join table is:

A	B	C
1	2	1
1	2	2
2	2	1
2	2	2
3	3	2

iv. We can clearly see that C is resolving to 1 & 2 for $A=1$.

v. We can also see that the number of rows is more in this table, compared to the original table R .

vi. The extra tuples are called *Spurious Tuples*.

• Example 1 (Lossless Decomposition):

◦ Table R :

A	B	C
1	2	1
2	2	2
3	3	2

- Table R_1 :

A	B
1	2
2	2
3	3

- Table R_2 :

A	C
1	1
2	2
3	2

- As we see, there are no `Spurious Tuples` when using `A` as the common attribute, following the above rule.